

Final Project Submission

Please fill out:

- Student name: Group 14 (Richard Gachiri, David Githaiga, Linah Ogumbeh, Allan Eshiter and Lemiso Eric)
- Student pace: part time
- Scheduled project review date/time:
- Instructor name:
- Blog post URL:

a) Introduction

This study aims to address the importance of identifying the factors that significantly influence home costs. Because conventional methods rely on anecdotal evidence or limited research, they often yield incorrect findings. This knowledge gap makes it more difficult for stakeholders to predict and assess changes in house prices. To solve this issue, we are using multiple regression modeling techniques to thoroughly analyze home sales data. Multiple regression is a useful tool for analyzing the relationship between different attributes and home sales prices since it takes into consideration the combined impact of several independent variables.

b) Business Problem

Real estate valuation poses significant challenges for Rittenhouse Brothers, particularly when determining property values influenced by unique features and the impact of renovations or upgrades. The subjective nature of valuation, varying interpretations among appraisers and professionals, and the absence of a purely objective methodology contribute to valuation discrepancies.

c) Main Objective

This Analysis aims to establish an objective property valuation model, focusing specifically on Property Unique features and the impact of renovations or upgrades. The primary goal is to minimize variations in valuation estimates and provide Rittenhouse Brothers with a more standardized and reliable method for assessing property values.

d) Specific Objectives

1. Perform exploratory data analysis to uncover connections between various variables and the target variable. This process aids in identifying pertinent variables for inclusion in a regression model.
2. Create a multiple regression model to forecast house sale prices by taking into account chosen independent variables and examining their influence on the dependent variable. Validate the model assumptions, evaluate its fitness for the data, and refine the model as needed.
3. Analyze the coefficients of the independent variables within the model to discern their individual contributions to house prices. Identify the most impactful factors infl

1. Loading Data

```
In [1]: # Your code here - remember to use markdown cells for comments as well!  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import scipy.stats as stats  
import seaborn as sns  
  
# Load data into pandas and engineer "make" feature  
data = pd.read_csv("./data/kc_house_data.csv")  
data.tail(50)
```

Out[1]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sq
21547	9406530090	10/20/2014	337000.0	4	2.50	2470	5100	2.0	NO	NONE	...	8 Good	2470	
21548	7168100015	10/9/2014	579950.0	5	2.75	3080	5752	2.0	NO	NONE	...	9 Better	3080	
21549	5007500120	2/26/2015	341780.0	4	2.75	2260	4440	2.0	NO	NONE	...	7 Average	2260	
21550	3528900770	4/23/2015	710200.0	4	3.00	1670	2642	2.0	NaN	NONE	...	8 Good	1350	
21551	9521100031	6/18/2014	690000.0	3	3.25	1540	1428	3.0	NO	NONE	...	9 Better	1540	
21552	524059330	1/30/2015	1700000.0	4	3.50	3830	8963	2.0	NO	NONE	...	10 Very Good	3120	
21553	6021503705	10/15/2014	329000.0	2	2.50	980	1020	3.0	NO	NONE	...	8 Good	980	
21554	3438501862	5/13/2014	330000.0	3	2.50	1450	5008	1.0	NO	NONE	...	7 Average	840	
21555	3345700207	5/2/2015	608500.0	4	3.50	2850	5577	2.0	NO	NONE	...	8 Good	1950	
21556	6056111067	7/7/2014	230000.0	3	1.75	1140	1201	2.0	NO	NONE	...	8 Good	1140	
21557	8562790760	5/20/2014	785000.0	4	3.50	3070	4684	2.0	NO	NONE	...	10 Very Good	2190	
21558	1931300090	5/7/2014	610950.0	3	3.00	1680	1570	3.0	NO	NONE	...	8 Good	1680	
21559	9578500790	11/11/2014	399950.0	3	2.50	3087	5002	2.0	NO	NONE	...	8 Good	3087	
21560	9253900271	1/7/2015	3570000.0	5	4.50	4850	10584	2.0	YES	EXCELLENT	...	10 Very Good	3540	
21561	3881900317	1/23/2015	579000.0	4	3.25	1900	2631	2.0	NO	NONE	...	9 Better	1250	
21562	567000385	6/23/2014	362500.0	2	1.50	940	1768	2.0	NaN	NONE	...	7 Average	940	
21563	7011201004	5/29/2014	645000.0	3	3.25	1730	1229	2.0	NO	AVERAGE	...	9 Better	1320	
21564	7853420110	10/3/2014	594866.0	3	3.00	2780	6000	2.0	NO	NONE	...	9 Better	2780	
21565	7853420110	5/4/2015	625000.0	3	3.00	2780	6000	2.0	NO	NONE	...	9 Better	2780	
21566	3052700432	11/12/2014	490000.0	3	2.25	1500	1290	2.0	NO	NONE	...	8 Good	1220	
21567	2025049203	6/10/2014	399950.0	2	1.00	710	1157	2.0	NaN	NONE	...	7 Average	710	

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sq
21568	952006823	12/2/2014	380000.0	3	2.50	1260	900	2.0	NO	NONE	...	7 Average	940	
21569	3832050760	8/28/2014	270000.0	3	2.50	1870	5000	2.0	NO	NONE	...	7 Average	1870	
21570	2767604724	10/15/2014	505000.0	2	2.50	1430	1201	3.0	NO	NONE	...	8 Good	1430	
21571	6632300207	3/5/2015	385000.0	3	2.50	1520	1488	3.0	NO	NONE	...	8 Good	1520	
21572	2767600688	11/13/2014	414500.0	2	1.50	1210	1278	2.0	NO	NONE	...	8 Good	1020	
21573	7570050450	9/10/2014	347500.0	3	2.50	2540	4760	2.0	NO	NONE	...	8 Good	2540	
21574	7430200100	5/14/2014	1220000.0	4	3.50	4910	9444	1.5	NO	NONE	...	11 Excellent	3110	
21575	4140940150	10/2/2014	572000.0	4	2.75	2770	3852	2.0	NO	NONE	...	8 Good	2770	
21576	1931300412	4/16/2015	475000.0	3	2.25	1190	1200	3.0	NO	NONE	...	8 Good	1190	
21577	8672200110	3/17/2015	1090000.0	5	3.75	4170	8142	2.0	NO	AVERAGE	...	10 Very Good	4170	
21578	5087900040	10/17/2014	350000.0	4	2.75	2500	5995	2.0	NaN	NONE	...	8 Good	2500	
21579	1972201967	10/31/2014	520000.0	2	2.25	1530	981	3.0	NO	NONE	...	8 Good	1480	
21580	7502800100	8/13/2014	679950.0	5	2.75	3600	9437	2.0	NO	NONE	...	9 Better	3600	
21581	191100405	4/21/2015	1580000.0	4	3.25	3410	10125	2.0	NO	NONE	...	10 Very Good	3410	
21582	8956200760	10/13/2014	541800.0	4	2.50	3118	7866	2.0	NaN	AVERAGE	...	9 Better	3118	
21583	7202300110	9/15/2014	810000.0	4	3.00	3990	7838	2.0	NO	NONE	...	9 Better	3990	
21584	249000205	10/15/2014	1540000.0	5	3.75	4470	8088	2.0	NO	NONE	...	11 Excellent	4470	
21585	5100403806	4/7/2015	467000.0	3	2.50	1425	1179	3.0	NO	NONE	...	8 Good	1425	
21586	844000965	6/26/2014	224000.0	3	1.75	1500	11968	1.0	NaN	NONE	...	6 Low Average	1500	
21587	7852140040	8/25/2014	507250.0	3	2.50	2270	5536	2.0	NaN	NONE	...	8 Good	2270	
21588	9834201367	1/26/2015	429000.0	3	2.00	1490	1126	3.0	NO	NONE	...	8 Good	1490	
21589	3448900210	10/14/2014	610685.0	4	2.50	2520	6023	2.0	NO	NaN	...	9 Better	2520	

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sq
21590	7936000429	3/26/2015	1010000.0	4	3.50	3510	7200	2.0	NO	NONE	...	9 Better	2600	
21591	2997800021	2/19/2015	475000.0	3	2.50	1310	1294	2.0	NO	NONE	...	8 Good	1180	
21592	263000018	5/21/2014	360000.0	3	2.50	1530	1131	3.0	NO	NONE	...	8 Good	1530	
21593	6600060120	2/23/2015	400000.0	4	2.50	2310	5813	2.0	NO	NONE	...	8 Good	2310	
21594	1523300141	6/23/2014	402101.0	2	0.75	1020	1350	2.0	NO	NONE	...	7 Average	1020	
21595	291310100	1/16/2015	400000.0	3	2.50	1600	2388	2.0	NaN	NONE	...	8 Good	1600	
21596	1523300157	10/15/2014	325000.0	2	0.75	1020	1076	2.0	NO	NONE	...	7 Average	1020	

50 rows × 21 columns

2. Exploratory Data

In [2]: data.shape

Out[2]: (21597, 21)

```
In [3]: data.describe() # Statistical distribution of the dataset
```

Out[3]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	yr_built	yr_renovated
count	2.159700e+04	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04	21597.000000	21597.000000	21597.000000	17755.000000
mean	4.580474e+09	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04	1.494096	1788.596842	1970.999676	83.636700
std	2.876736e+09	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04	0.539683	827.759761	29.375234	399.946400
min	1.000102e+06	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02	1.000000	370.000000	1900.000000	0.000000
25%	2.123049e+09	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03	1.000000	1190.000000	1951.000000	0.000000
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03	1.500000	1560.000000	1975.000000	0.000000
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04	2.000000	2210.000000	1997.000000	0.000000
max	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	9410.000000	2015.000000	2015.000000

```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   id                    21597 non-null  int64  
 1   date                  21597 non-null  object  
 2   price                 21597 non-null  float64 
 3   bedrooms              21597 non-null  int64  
 4   bathrooms              21597 non-null  float64 
 5   sqft_living           21597 non-null  int64  
 6   sqft_lot              21597 non-null  int64  
 7   floors                 21597 non-null  float64 
 8   waterfront            19221 non-null  object  
 9   view                  21534 non-null  object  
10   condition              21597 non-null  object  
11   grade                 21597 non-null  object  
12   sqft_above            21597 non-null  int64  
13   sqft_basement         21597 non-null  object  
14   yr_built              21597 non-null  int64  
15   yr_renovated          17755 non-null  float64 
16   zipcode               21597 non-null  int64  
17   lat                   21597 non-null  float64 
18   long                  21597 non-null  float64 
19   sqft_living15         21597 non-null  int64  
20   sqft_lot15            21597 non-null  int64  
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB
```

```
In [5]: data.columns
```

```
Out[5]: Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',
              'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',
              'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',
              'lat', 'long', 'sqft_living15', 'sqft_lot15'],
              dtype='object')
```



```
In [6]: data.dtypes
```

```
Out[6]: id                int64  
date                object  
price              float64  
bedrooms           int64  
bathrooms          float64  
sqft_living         int64  
sqft_lot            int64  
floors             float64  
waterfront         object  
view               object  
condition          object  
grade              object  
sqft_above          int64  
sqft_basement       object  
yr_built            int64  
yr_renovated        float64  
zipcode             int64  
lat                 float64  
long                float64  
sqft_living15       int64  
sqft_lot15          int64  
dtype: object
```

3. Data Preparation

a) Data Cleaning

```
In [7]: data.isnull().sum()
```

```
Out[7]: id                0  
date                0  
price              0  
bedrooms           0  
bathrooms          0  
sqft_living         0  
sqft_lot            0  
floors              0  
waterfront         2376  
view                63  
condition           0  
grade               0  
sqft_above          0  
sqft_basement       0  
yr_built            0  
yr_renovated        3842  
zipcode             0  
lat                 0  
long                0  
sqft_living15       0  
sqft_lot15          0  
dtype: int64
```

```
In [8]: missing_values = data.isna() # Checks missing values
missing_values.head()
```

Out[8]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	yr_built	yr_renovated
0	False	False	False	False	False	False	False	False	True	False	...	False	False	False	False	False
1	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	...	False	False	False	False	False

5 rows × 21 columns



```
In [9]: def missing_values(data):
# identify the total missing values per column
# sort in order
miss = data.isnull().sum().sort_values(ascending = False)

# calculate percentage of the missing values
percentage_miss = (data.isnull().sum() / len(data)).sort_values(ascending = False)

# store in a dataframe
missing = pd.DataFrame({"Missing Values": miss, "Percentage": percentage_miss}).reset_index()

# remove values that are missing
missing.drop(missing[missing["Percentage"] == 0].index, inplace = True)

return missing
```

```
In [10]: missing_data = missing_values(data)
missing_data
```

Out[10]:

	index	Missing Values	Percentage
0	yr_renovated	3842	0.177895
1	waterfront	2376	0.110015
2	view	63	0.002917

Replace the rows with the missing values for **View** with the mode of view

```
In [11]: data['view'].value_counts()
```

Out[11]:

NONE	19422
AVERAGE	957
GOOD	508
FAIR	330
EXCELLENT	317

Name: view, dtype: int64

```
In [12]: data['view'] = data['view'].fillna(data['view'].mode().iloc[0])
```

```
In [13]: data.isna().sum()
```

```
Out[13]: id                0
         date              0
         price             0
         bedrooms          0
         bathrooms         0
         sqft_living        0
         sqft_lot          0
         floors            0
         waterfront        2376
         view              0
         condition         0
         grade             0
         sqft_above         0
         sqft_basement      0
         yr_built           0
         yr_renovated       3842
         zipcode           0
         lat               0
         long              0
         sqft_living15      0
         sqft_lot15         0
         dtype: int64
```

Drop the rows with missing values in **Waterfront**

```
In [14]: data['waterfront'].value_counts()
```

```
Out[14]: NO      19075
         YES       146
         Name: waterfront, dtype: int64
```

```
In [15]: # Mode imputation for the missing values in waterfront
         data['waterfront'] = data['waterfront'].fillna(data['waterfront'].mode().iloc[0])
```

```
In [16]: # check for the count unique values in the year renovated column  
data['yr_renovated'].value_counts()
```

```
Out[16]: 0.0      17011  
         2014.0     73  
         2013.0     31  
         2003.0     31  
         2007.0     30  
         ...  
         1951.0      1  
         1953.0      1  
         1946.0      1  
         1976.0      1  
         1948.0      1  
Name: yr_renovated, Length: 70, dtype: int64
```

```
In [17]: # Replace null with the most frequent value  
data['yr_renovated'].fillna(0,inplace=True)
```

```
In [18]: data.isna().sum()
```

```
Out[18]: id                0
         date              0
         price             0
         bedrooms          0
         bathrooms         0
         sqft_living        0
         sqft_lot           0
         floors             0
         waterfront        0
         view              0
         condition         0
         grade             0
         sqft_above         0
         sqft_basement      0
         yr_built           0
         yr_renovated       0
         zipcode            0
         lat                0
         long               0
         sqft_living15      0
         sqft_lot15         0
         dtype: int64
```

No Missing Values

```
In [19]: # Drop rows with '?' in 'sqft_basement'
         data = data[data['sqft_basement'] != '?']

         # Reset the index of the DataFrame
         data.reset_index(drop=True, inplace=True)
```

```
In [20]: # Convert 'sqft_basement' column to float
         data['sqft_basement'] = data['sqft_basement'].astype(float)
```

```
In [21]: ## Covert the data type for data  
data['date'] = pd.to_datetime(data['date'], format = '%m/%d/%Y' )
```

```
In [22]: data.shape
```

```
Out[22]: (21143, 21)
```

b) Encoding Categorical Columns in the Data

```
In [23]: # Encode the column waterfront  
waterfront_category_order = {  
    'NO': 0,  
    "YES" : 1  
}  
  
# Perform ordinal encoding on the "waterfront" column  
data["waterfront_encoded"] = data["waterfront"].map(waterfront_category_order)  
  
# Replace the original "waterfront" column with the encoded values  
data["waterfront"] = data["waterfront_encoded"]  
  
data['waterfront'].value_counts()
```

```
Out[23]: 0    21001  
        1      142  
        Name: waterfront, dtype: int64
```

```
In [24]: # Let's check the count of the encoded column of waterfront  
data['waterfront'].value_counts()
```

```
Out[24]: 0    21001  
        1      142  
        Name: waterfront, dtype: int64
```



```
In [25]: # Encode the column view
# Define the order of the categories
view_category_order = {
    "NONE": 0,
    "AVERAGE": 1,
    "GOOD": 2,
    "FAIR": 3,
    "EXCELLENT": 4
}

# Perform ordinal encoding on the "view" column
data["view_encoded"] = data["view"].map(view_category_order)

# Replace the original "view" column with the encoded values
data["view"] = data["view_encoded"]

data['view'].value_counts()
```

```
Out[25]: 0    19079
         1     930
         2     496
         3     327
         4     311
         Name: view, dtype: int64
```

```
In [26]: # Let's encode the condition column
condition_category_order = {
    "Poor": 1,
    "Fair": 2,
    "Average": 3,
    "Good": 4,
    "Very Good": 5
}

# Perform ordinal encoding on the "condition" column
data["condition_encoded"] = data["condition"].map(condition_category_order)

# Replace the original "condition" column with the encoded values
data["condition"] = data["condition_encoded"]

data['condition'].value_counts()
```

```
Out[26]: 3    13726
         4     5557
         5     1666
         2      166
         1        28
         Name: condition, dtype: int64
```

```
In [27]: # Encode the grade column
grade_category_order = {
    "3 Poor": 3,
    "5 Fair": 5,
    "4 Low": 4,
    "6 Low Average": 6,
    "7 Average": 7,
    "8 Good": 8,
    "9 Better": 9,
    "10 Very Good": 10,
    "11 Excellent": 11,
    "12 Luxury": 12,
    "13 Mansion": 13
}

# Perform ordinal encoding on the "grade" column
data["grade_encoded"] = data["grade"].replace(grade_category_order)

# Replace the original "grade" column with the encoded values
data["grade"] = data["grade_encoded"]

data['grade'].value_counts()
```

```
Out[27]: 7      8788
        8      5933
        9      2557
        6      1997
       10      1112
       11       391
        5       235
       12        89
        4        27
       13        13
        3         1
        Name: grade, dtype: int64
```

```
In [28]: encoded_data = data
encoded_data
```

Out[28]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	yr_renovated	zipcode	lat	lc
0	7129300520	2014-10-13	221900.0	3	1.00	1180	5650	1.0	0	0	...	0.0	98178	47.5112	-122.1
1	6414100192	2014-12-09	538000.0	3	2.25	2570	7242	2.0	0	0	...	1991.0	98125	47.7210	-122.1
2	5631500400	2015-02-25	180000.0	2	1.00	770	10000	1.0	0	0	...	0.0	98028	47.7379	-122.1
3	2487200875	2014-12-09	604000.0	4	3.00	1960	5000	1.0	0	0	...	0.0	98136	47.5208	-122.1
4	1954400510	2015-02-18	510000.0	3	2.00	1680	8080	1.0	0	0	...	0.0	98074	47.6168	-122.0
...
21138	263000018	2014-05-21	360000.0	3	2.50	1530	1131	3.0	0	0	...	0.0	98103	47.6993	-122.1
21139	6600060120	2015-02-23	400000.0	4	2.50	2310	5813	2.0	0	0	...	0.0	98146	47.5107	-122.1
21140	1523300141	2014-06-23	402101.0	2	0.75	1020	1350	2.0	0	0	...	0.0	98144	47.5944	-122.1
21141	291310100	2015-01-16	400000.0	3	2.50	1600	2388	2.0	0	0	...	0.0	98027	47.5345	-122.0
21142	1523300157	2014-10-15	325000.0	2	0.75	1020	1076	2.0	0	0	...	0.0	98144	47.5941	-122.1

21143 rows × 25 columns



In [29]: data.dtypes

```
Out[29]: id                int64
         date              datetime64[ns]
         price             float64
         bedrooms          int64
         bathrooms         float64
         sqft_living        int64
         sqft_lot           int64
         floors             float64
         waterfront        int64
         view              int64
         condition         int64
         grade             int64
         sqft_above         int64
         sqft_basement      float64
         yr_built           int64
         yr_renovated       float64
         zipcode           int64
         lat               float64
         long              float64
         sqft_living15      int64
         sqft_lot15        int64
         waterfront_encoded int64
         view_encoded       int64
         condition_encoded  int64
         grade_encoded      int64
         dtype: object
```

```
In [30]: # Drop the encoded columns
data_updated= encoded_data.drop(["waterfront_encoded", "view_encoded", "condition_encoded", "grade_encoded"], axis=1)

# Check the updated dataset
data_updated
```

Out[30]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	...	grade	sqft_above	sqft_basement	yr.
0	7129300520	2014-10-13	221900.0	3	1.00	1180	5650	1.0	0	0	...	7	1180	0.0	
1	6414100192	2014-12-09	538000.0	3	2.25	2570	7242	2.0	0	0	...	7	2170	400.0	
2	5631500400	2015-02-25	180000.0	2	1.00	770	10000	1.0	0	0	...	6	770	0.0	
3	2487200875	2014-12-09	604000.0	4	3.00	1960	5000	1.0	0	0	...	7	1050	910.0	
4	1954400510	2015-02-18	510000.0	3	2.00	1680	8080	1.0	0	0	...	8	1680	0.0	
...	
21138	263000018	2014-05-21	360000.0	3	2.50	1530	1131	3.0	0	0	...	8	1530	0.0	
21139	6600060120	2015-02-23	400000.0	4	2.50	2310	5813	2.0	0	0	...	8	2310	0.0	
21140	1523300141	2014-06-23	402101.0	2	0.75	1020	1350	2.0	0	0	...	7	1020	0.0	
21141	291310100	2015-01-16	400000.0	3	2.50	1600	2388	2.0	0	0	...	8	1600	0.0	
21142	1523300157	2014-10-15	325000.0	2	0.75	1020	1076	2.0	0	0	...	7	1020	0.0	

21143 rows × 21 columns



```
In [31]: # New feature (number of years from yr_built to yr_renovated)
data_updated['YrsB4Renovation'] = data_updated['yr_renovated'] - data_updated['yr_built']
```


```
In [32]: ### Drop the unnecessary
data_updated.drop(['id', 'yr_built', 'yr_renovated'], axis=1, inplace=True)
data_updated.head()
```

Out[32]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	zipcode
0	2014-10-13	221900.0	3	1.00	1180	5650	1.0	0	0	3	7	1180	0.0	98178 47.
1	2014-12-09	538000.0	3	2.25	2570	7242	2.0	0	0	3	7	2170	400.0	98125 47.
2	2015-02-25	180000.0	2	1.00	770	10000	1.0	0	0	3	6	770	0.0	98028 47.
3	2014-12-09	604000.0	4	3.00	1960	5000	1.0	0	0	5	7	1050	910.0	98136 47.
4	2015-02-18	510000.0	3	2.00	1680	8080	1.0	0	0	3	8	1680	0.0	98074 47.

4. Explanatory Data Analysis

```
In [33]: # Distribution of the target variable
plt.figure(figsize=(8, 6))
sns.histplot(data_updated, x='price', kde=True)
plt.title("Distribution of Price")
plt.show()
```

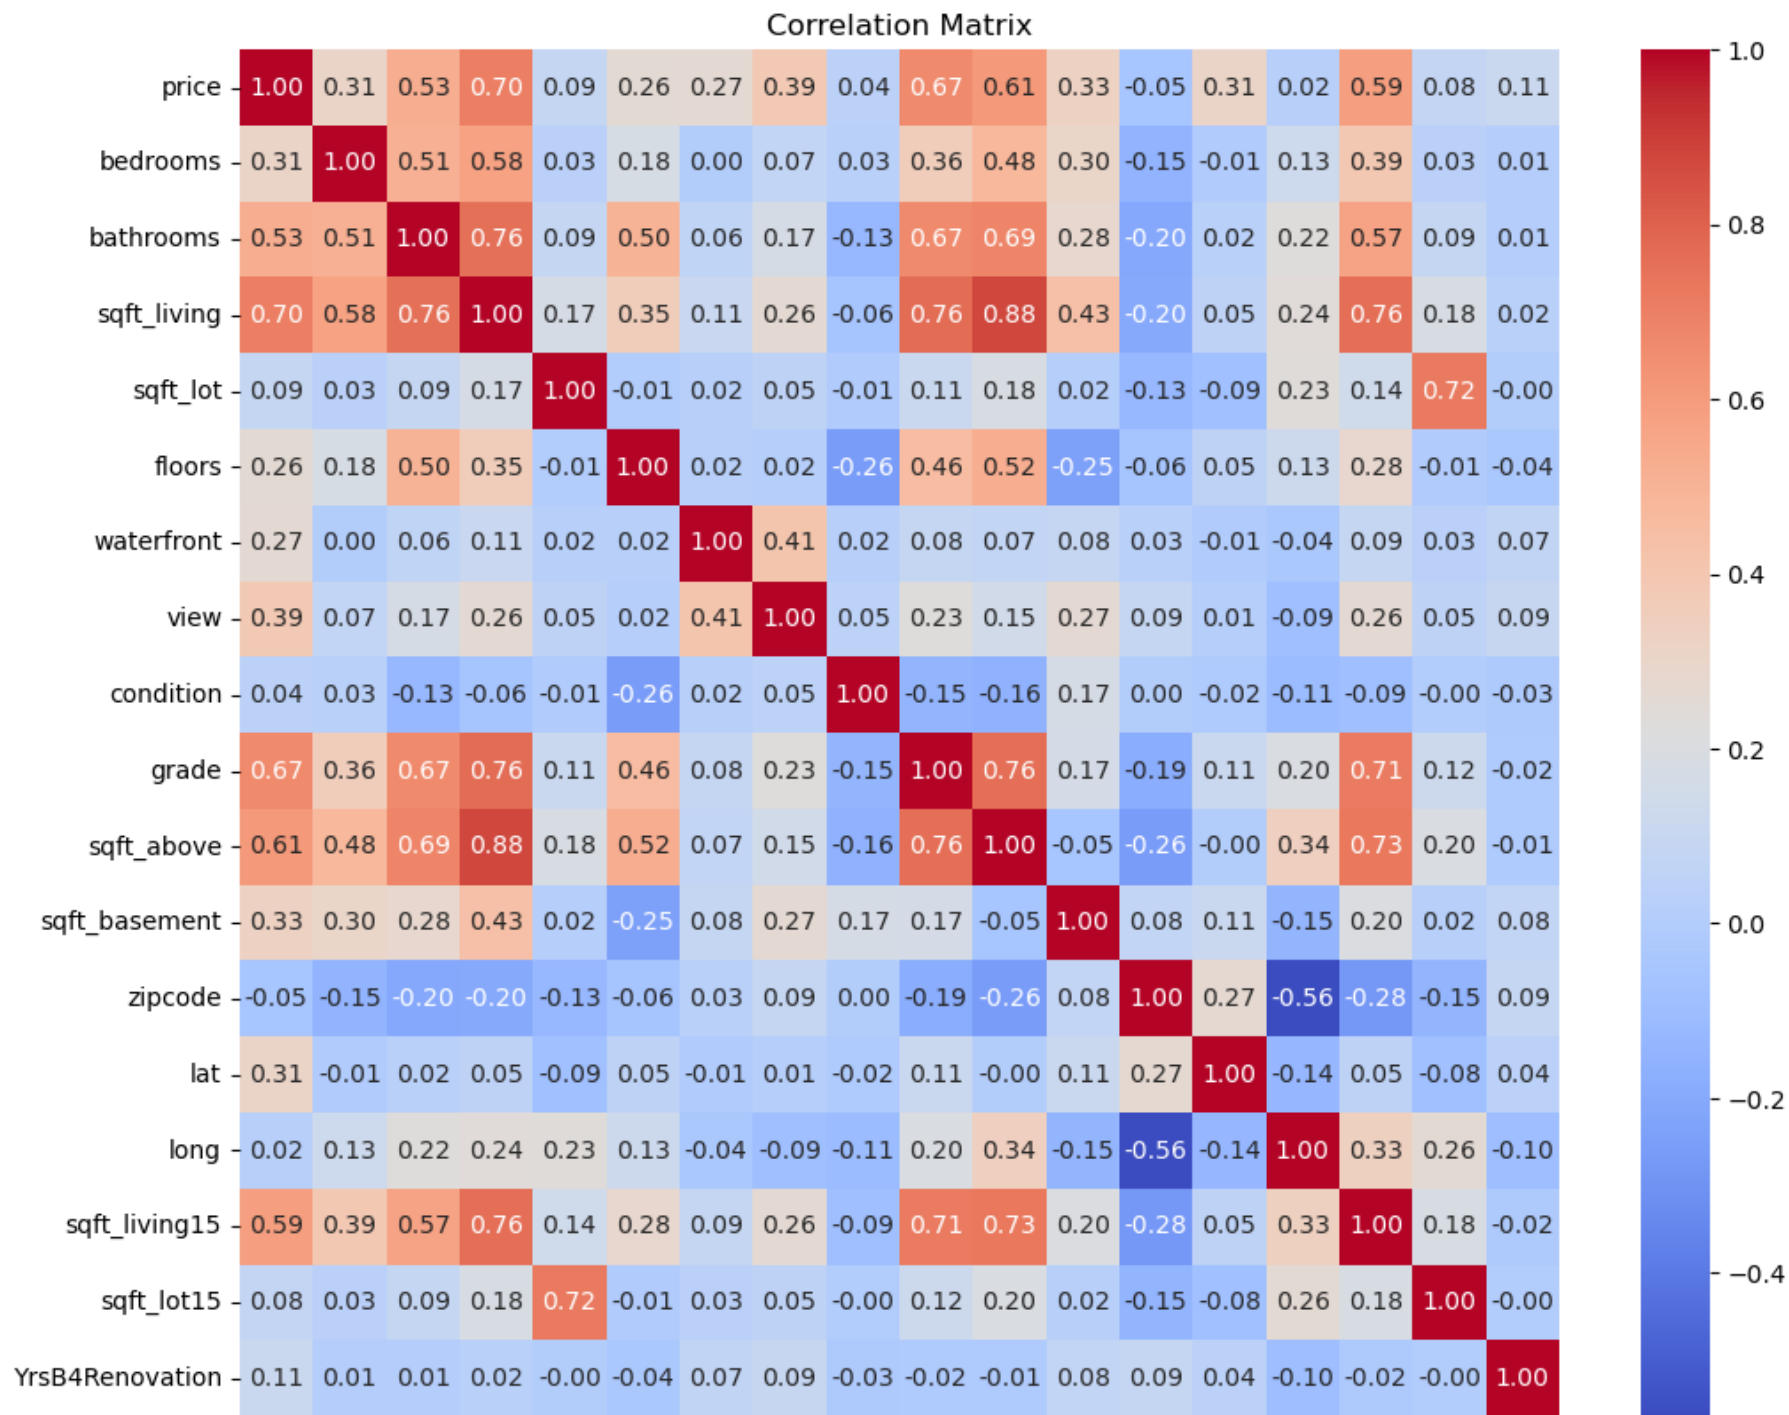
The diagram shows a positively skewed distribution of the price

```
In [34]: ### Check Correlation
data_updated.corr(numeric_only=True)
```

Out[34]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_b
price	1.000000	0.309204	0.525889	0.702328	0.087937	0.256355	0.265969	0.387878	0.035290	0.667738	0.605143	
bedrooms	0.309204	1.000000	0.513840	0.577998	0.032262	0.179044	0.000061	0.072502	0.025710	0.356882	0.479210	
bathrooms	0.525889	0.513840	1.000000	0.755278	0.087884	0.504071	0.064626	0.171079	-0.128015	0.666493	0.686456	
sqft_living	0.702328	0.577998	0.755278	1.000000	0.172941	0.354688	0.106039	0.263900	-0.061294	0.763101	0.876678	
sqft_lot	0.087937	0.032262	0.087884	0.172941	1.000000	-0.007522	0.021211	0.050720	-0.009671	0.113402	0.183461	
floors	0.256355	0.179044	0.504071	0.354688	-0.007522	1.000000	0.020307	0.016255	-0.264280	0.459214	0.524225	-
waterfront	0.265969	0.000061	0.064626	0.106039	0.021211	0.020307	1.000000	0.408943	0.016728	0.083602	0.072201	
view	0.387878	0.072502	0.171079	0.263900	0.050720	0.016255	0.408943	1.000000	0.046585	0.227733	0.150351	
condition	0.035290	0.025710	-0.128015	-0.061294	-0.009671	-0.264280	0.016728	0.046585	1.000000	-0.148508	-0.159289	
grade	0.667738	0.356882	0.666493	0.763101	0.113402	0.459214	0.083602	0.227733	-0.148508	1.000000	0.756382	
sqft_above	0.605143	0.479210	0.686456	0.876678	0.183461	0.524225	0.072201	0.150351	-0.159289	0.756382	1.000000	-
sqft_basement	0.325008	0.302683	0.282693	0.434576	0.015533	-0.245144	0.084949	0.266334	0.170972	0.168023	-0.052293	
zipcode	-0.053166	-0.152628	-0.204306	-0.198906	-0.129355	-0.058976	0.028152	0.085423	0.003437	-0.185131	-0.260814	
lat	0.306507	-0.009521	0.024994	0.052986	-0.085457	0.048977	-0.011567	0.014825	-0.015592	0.113380	-0.000346	
long	0.022101	0.131093	0.224479	0.240797	0.230583	0.125922	-0.036371	-0.089194	-0.106285	0.199187	0.344543	-
sqft_living15	0.586415	0.392476	0.570129	0.756389	0.143428	0.280294	0.085174	0.261981	-0.093493	0.713574	0.731016	
sqft_lot15	0.083192	0.030479	0.088834	0.184466	0.720649	-0.011545	0.030250	0.052490	-0.004860	0.121921	0.195638	
YrsB4Renovation	0.110193	0.005652	0.006731	0.023686	-0.000082	-0.035371	0.070794	0.093260	-0.026427	-0.019954	-0.014552	

```
In [35]: # Correlation matrix
plt.figure(figsize=(12, 10))
sns.heatmap(data_updated.corr(numeric_only=True), annot=True, cmap='coolwarm', fmt='.2f')
plt.title("Correlation Matrix")
plt.show()
```

price
bedrooms
bathrooms
sqft_living
sqft_lot
floors
waterfront
view
condition
grade
sqft_above
sqft_basement
zipcode
lat
long
sqft_living15
sqft_lot15
YrsB4Renovation

We have so many numbers on the heatmap but we are only interested in the large values that are equal or more than 0.75

This will check for multicollinearity within independent variables. High multicollinearity might lead to poor performance of our model

```
In [36]: # Scatter plots
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
sns.scatterplot(data_updated, x='sqft_living', y='price')
plt.title("Price vs. Sqft Living")


plt.subplot(1, 2, 2)
sns.scatterplot(data_updated, x='grade', y='price')
plt.title("Price vs. Grade")


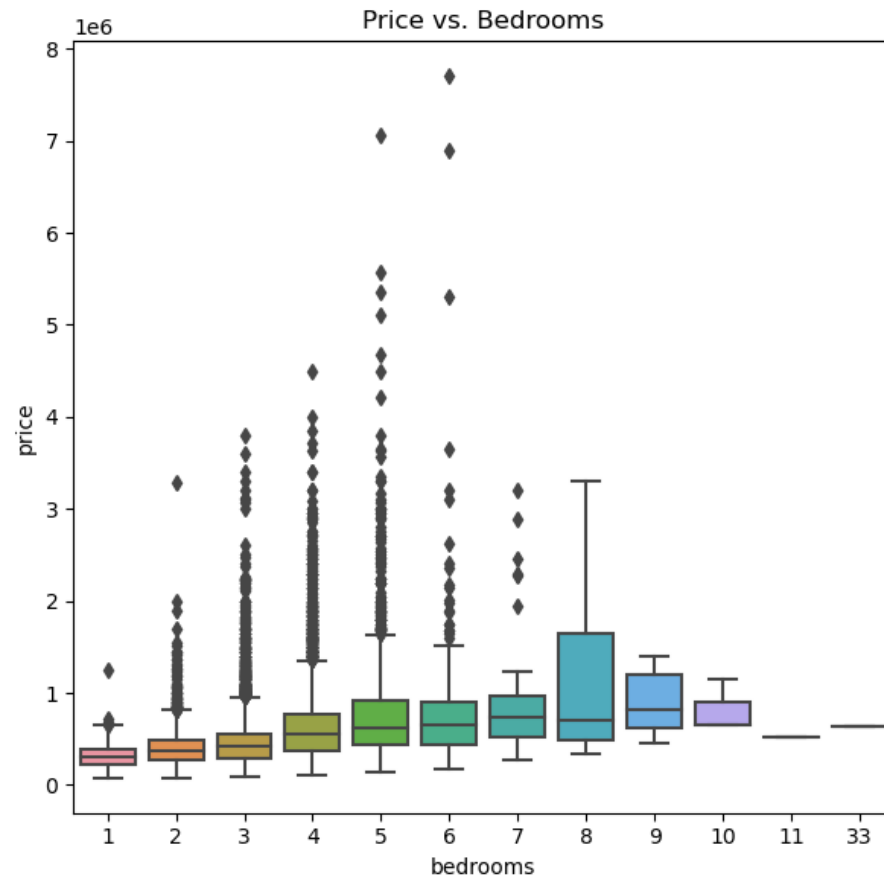
plt.tight_layout()
plt.show()

# Box plots
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
sns.boxplot(data_updated, x='bedrooms', y='price')
plt.title("Price vs. Bedrooms")

plt.subplot(1, 2, 2)
sns.boxplot(data_updated, x='floors', y='price')
plt.title("Price vs. Floors")

plt.tight_layout()
plt.show()
```






```
In [37]: data = data_updated  
data.head()
```

Out[37]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	condition	grade	sqft_above	sqft_basement	zipcode	
0	2014-10-13	221900.0	3	1.00	1180	5650	1.0	0	0	3	7	1180	0.0	98178	47.
1	2014-12-09	538000.0	3	2.25	2570	7242	2.0	0	0	3	7	2170	400.0	98125	47.
2	2015-02-25	180000.0	2	1.00	770	10000	1.0	0	0	3	6	770	0.0	98028	47.
3	2014-12-09	604000.0	4	3.00	1960	5000	1.0	0	0	5	7	1050	910.0	98136	47.
4	2015-02-18	510000.0	3	2.00	1680	8080	1.0	0	0	3	8	1680	0.0	98074	47.



```
In [38]: data.shape
```

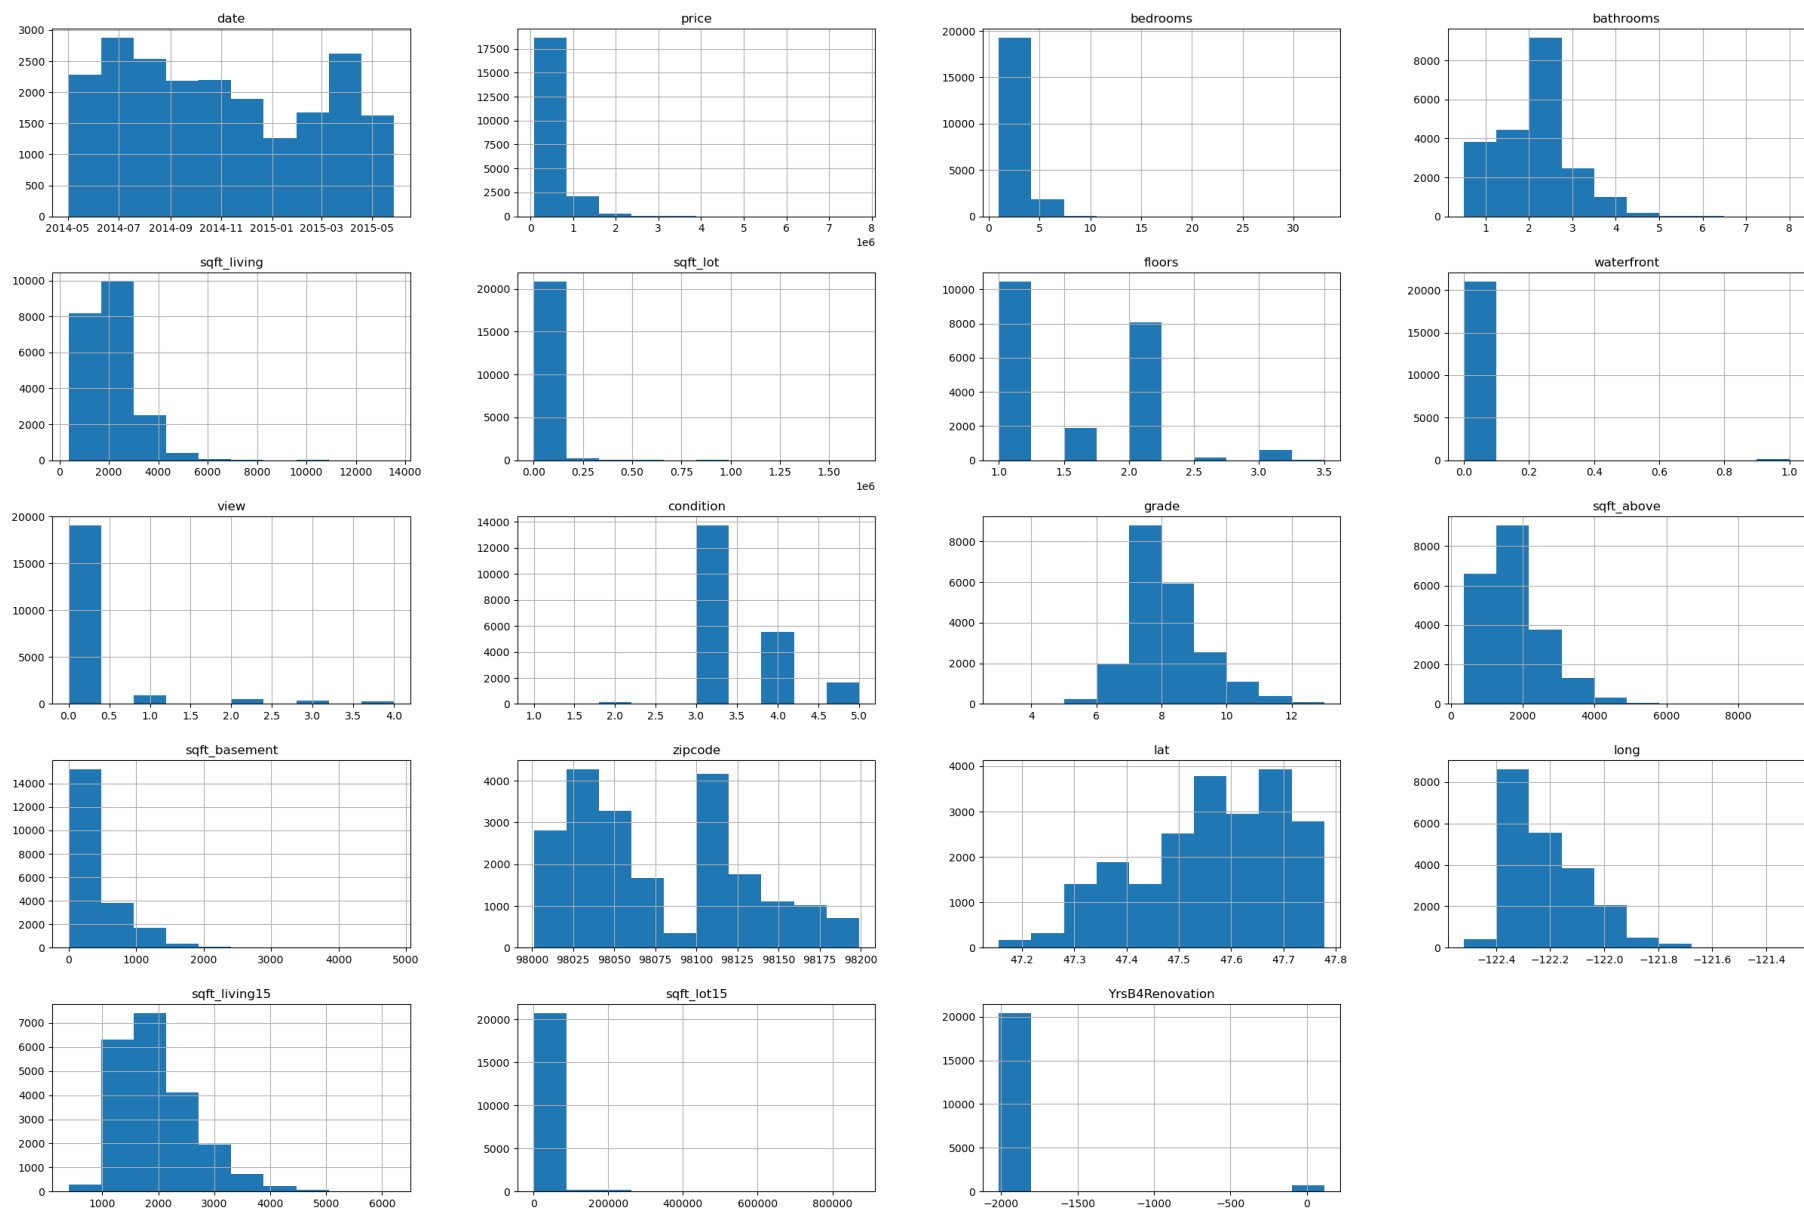
Out[38]: (21143, 19)

5. Model Iterations

```
In [39]: data.corr(numeric_only=True)['price']
```

```
Out[39]: price            1.000000  
bedrooms          0.309204  
bathrooms          0.525889  
sqft_living        0.702328  
sqft_lot           0.087937  
floors             0.256355  
waterfront         0.265969  
view               0.387878  
condition          0.035290  
grade              0.667738  
sqft_above         0.605143  
sqft_basement      0.325008  
zipcode            -0.053166  
lat                0.306507  
long               0.022101  
sqft_living15      0.586415  
sqft_lot15         0.083192  
YrsB4Renovation    0.110193  
Name: price, dtype: float64
```

```
In [40]: data.hist(figsize=(30,20))  
plt.show()
```



a) Basic Model: Simple Linear Regression with 'sqft_living': (Baseline Model)

```
In [41]: import statsmodels.api as sm
```

```
In [42]: # Selecting the feature and target variable
X_baseline = data[['sqft_living']] # Feature (independent variable)
y = data['price'] # Target (dependent variable)

baseline_model = sm.OLS(y, sm.add_constant(X_baseline))
baseline_results = baseline_model.fit()

print(baseline_results.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:                0.493
Model:                  OLS       Adj. R-squared:           0.493
Method:                 Least Squares   F-statistic:         2.058e+04
Date:                  Tue, 02 Jan 2024   Prob (F-statistic):   0.00
Time:                  19:53:21    Log-Likelihood:      -2.9378e+05
No. Observations:      21143      AIC:                 5.876e+05
Df Residuals:          21141      BIC:                 5.876e+05
Df Model:               1
Covariance Type:       nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
const          -4.513e+04    4462.443    -10.113    0.000    -5.39e+04    -3.64e+04
sqft_living     281.4327        1.962     143.454    0.000     277.587     285.278
=====
Omnibus:            14518.924   Durbin-Watson:           1.986
Prob(Omnibus):      0.000   Jarque-Bera (JB):       535834.159
Skew:               2.825   Prob(JB):               0.00
Kurtosis:           27.006   Cond. No.               5.63e+03
=====
```


Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.63e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Regression Line:

```
In [43]: fig, ax = plt.subplots(1,2, figsize= (10,5))
data.plot.scatter(x="sqft_living", y="price", label="Data points", ax=ax[0])
sm.graphics.abline_plot(model_results=baseline_results, label="Regression line", ax=ax[0], color="black")
ax[0].legend();

ax[1].scatter(data["price"], baseline_results.resid)
ax[1].axhline(y=0, color="black")
ax[1].set_xlabel("sqft_living")
ax[1].set_ylabel("residuals");
```



The residuals:

```
In [44]: # Log-transform target and predictor variables in baseline model
Xlog_baseline = np.log(data[['sqft_living']]) # Feature (independent variable)
ylog = np.log(data['price']) # Target (dependent variable)

log_baseline_model = sm.OLS(ylog, sm.add_constant(Xlog_baseline))
log_baseline_results = log_baseline_model.fit()

print(log_baseline_results.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          price    R-squared:                0.455
Model:                  OLS      Adj. R-squared:           0.455
Method:                 Least Squares    F-statistic:        1.766e+04
Date:                  Tue, 02 Jan 2024    Prob (F-statistic):    0.00
Time:                  19:53:22    Log-Likelihood:       -10023.
No. Observations:      21143    AIC:                  2.005e+04
Df Residuals:          21141    BIC:                  2.007e+04
Df Model:               1
Covariance Type:       nonrobust
=====
                        coef    std err          t      P>|t|      [0.025    0.975]
-----
const                6.7210      0.048    140.939      0.000      6.628      6.814
sqft_living          0.8380      0.006    132.895      0.000      0.826      0.850
=====
Omnibus:              120.771    Durbin-Watson:        1.979
Prob(Omnibus):         0.000    Jarque-Bera (JB):     111.909
Skew:                  0.144    Prob(JB):             5.00e-25
Kurtosis:              2.790    Cond. No.             137.
=====

```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

A 1% increase in sqft_living leads to a 0.84% increase in house price


```
In [45]: fig, ax = plt.subplots(1,2, figsize= (10,5))
ax[0].scatter(np.log(data['sqft_living']), np.log(data["price"]), label="Data points")
sm.graphics.abline_plot(model_results=log_baseline_results, label="Regression line", ax=ax[0], color="black")
ax[0].legend();

ax[1].scatter(np.log(data["price"]), log_baseline_results.resid)
ax[1].axhline(y=0, color="black")
ax[1].set_xlabel("sqft_living")
ax[1].set_ylabel("residuals");
```



Despite having a slightly lower R-squared value, the model with log-transformation baseline model performs better based on the relatively AIC and BIC values

b) Add More Features for Multiple Regression with raw data

```
In [46]: features = ['sqft_living', 'floors', 'lat', 'zipcode', 'view']

X_one = data[features].copy()
X_one['sqft_living'] = np.log(X_one['sqft_living'])
y = np.log(data['price'])

subset_model_one = sm.OLS(y, sm.add_constant(X_one))
subset_results_one = subset_model_one.fit()

print(subset_results_one.summary())
```

OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:                0.672
Model:                  OLS      Adj. R-squared:           0.672
Method:                 Least Squares    F-statistic:          8673.
Date:                  Tue, 02 Jan 2024    Prob (F-statistic):    0.00
Time:                  19:53:22    Log-Likelihood:       -4649.1
No. Observations:      21143    AIC:                  9310.
Df Residuals:          21137    BIC:                  9358.
Df Model:              5
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-33.6187	3.972	-8.464	0.000	-41.404	-25.834
sqft_living	0.7181	0.006	128.021	0.000	0.707	0.729
floors	0.0692	0.004	16.656	0.000	0.061	0.077
lat	1.6340	0.016	104.540	0.000	1.603	1.665
zipcode	-0.0004	4.19e-05	-8.910	0.000	-0.000	-0.000
view	0.1508	0.003	48.050	0.000	0.145	0.157

```

=====
Omnibus:                387.036    Durbin-Watson:          1.994
Prob(Omnibus):           0.000    Jarque-Bera (JB):       504.870
Skew:                    0.246    Prob(JB):               2.34e-110
Kurtosis:                3.576    Cond. No.               1.88e+08
=====

```

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.88e+08. This might indicate that there are strong multicollinearity or other numerical problems.

c) Add More Features for Multiple Regression and using standardised predictor variables

```
In [47]: X_two = X_one.copy()
for col in X_two.columns:
    X_two[col] = (X_two[col] - X_two[col].mean())/X_two[col].std()

subset_model_two = sm.OLS(y, sm.add_constant(X_two))
subset_results_two = subset_model_two.fit()

print(subset_results_two.summary())
```

OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:                0.672
Model:                  OLS      Adj. R-squared:           0.672
Method:                 Least Squares    F-statistic:          8673.
Date:                  Tue, 02 Jan 2024    Prob (F-statistic):    0.00
Time:                  19:53:22    Log-Likelihood:       -4649.1
No. Observations:      21143    AIC:                  9310.
Df Residuals:          21137    BIC:                  9358.
Df Model:               5
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	13.0484	0.002	6292.472	0.000	13.044	13.052
sqft_living	0.3045	0.002	128.021	0.000	0.300	0.309
floors	0.0373	0.002	16.656	0.000	0.033	0.042
lat	0.2265	0.002	104.540	0.000	0.222	0.231
zipcode	-0.0200	0.002	-8.910	0.000	-0.024	-0.016
view	0.1038	0.002	48.050	0.000	0.100	0.108

```

=====
Omnibus:                387.036    Durbin-Watson:          1.994
Prob(Omnibus):          0.000    Jarque-Bera (JB):       504.870
Skew:                   0.246    Prob(JB):               2.34e-110
Kurtosis:               3.576    Cond. No.                1.73
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Based on the standardised model, sqft_living has the largest impact on the house prices since it has the largest co-efficient.

```
In [48]: #Mean centered predictors
X_centered = X_one.copy()
for col in X_centered.columns:
    X_centered[col] -= X_centered[col].mean()

model = sm.OLS(y, sm.add_constant(X_centered)).fit()
print(model.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          price    R-squared:                0.672
Model:                  OLS      Adj. R-squared:           0.672
Method:                 Least Squares    F-statistic:        8673.
Date:                   Tue, 02 Jan 2024    Prob (F-statistic):    0.00
Time:                   19:53:22    Log-Likelihood:       -4649.1
No. Observations:      21143    AIC:                  9310.
Df Residuals:          21137    BIC:                  9358.
Df Model:               5
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	13.0484	0.002	6292.472	0.000	13.044	13.052
sqft_living	0.7181	0.006	128.021	0.000	0.707	0.729
floors	0.0692	0.004	16.656	0.000	0.061	0.077
lat	1.6340	0.016	104.540	0.000	1.603	1.665
zipcode	-0.0004	4.19e-05	-8.910	0.000	-0.000	-0.000
view	0.1508	0.003	48.050	0.000	0.145	0.157

```
=====
Omnibus:                 387.036    Durbin-Watson:           1.994
Prob(Omnibus):            0.000    Jarque-Bera (JB):        504.870
Skew:                     0.246    Prob(JB):                 2.34e-110
Kurtosis:                 3.576    Cond. No.                  404.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Since only price, i.e. the target variable, and sqft_living has been log-transformed, the model interpretation means that for 1 unit increase in floors leads to a 9.4% increase in house prices. A one level improvement in condition leads to a 7.34% increase in house prices while a one level improvement in view leads to a 14% increase house price

d) Mean Square Errors

```
In [49]: from sklearn.metrics import mean_squared_error
```

```
In [50]: # Make predictions using the baseline model
predictions = baseline_results.predict(sm.add_constant(X_baseline))

# Calculate MSE
mse = mean_squared_error(y, predictions)
print(f"Mean Squared Error: {mse}")
```

Mean Squared Error: 358961456587.6772

```
In [51]: # Make predictions using the log-transformed baseline model
predictions = log_baseline_results.predict(sm.add_constant(Xlog_baseline))

# Calculate MSE
mse = mean_squared_error(y, predictions)
print(f"Mean Squared Error: {mse}")
```

Mean Squared Error: 0.15111301799281243

```
In [52]: # Make predictions using the mean-centered model
predictions = subset_results_one.predict(sm.add_constant(X_one))

# Calculate MSE
mse = mean_squared_error(y, predictions)
print(f"Mean Squared Error: {mse}")
```

Mean Squared Error: 0.09089019618976422

```
In [53]: # Make predictions using the standardised model
predictions = subset_results_two.predict(sm.add_constant(X_two))

# Calculate MSE
mse = mean_squared_error(y, predictions)
print(f"Mean Squared Error: {mse}")
```

Mean Squared Error: 0.09089019618976431


6. Assumptions

We will inspect **linearity, independence, normality and Equal Variance**

a) The Linearity Assumption

```
In [54]: fig, ax = plt.subplots()

ax.scatter(y, subset_results_two.resid, color="green")
ax.axhline(y=0, color="black")
ax.set_xlabel("y")
ax.set_ylabel("residuals")
ax.set_title("Linear Relationship Residual Plot");
```



```
In [55]: # performing a rainbow test to test linearity statistically  
from statsmodels.stats.diagnostic import linear_rainbow  
linear_rainbow(subset_results_one)
```

```
Out[55]: (0.9604358952019966, 0.9810006700501094)
```

b) Independence Assumption


The Durbin-Watson value in our instance is 1.994, which is quite near to 2, indicating that the residuals show little indication of first-order autocorrelation. This is advantageous since it upholds the regression model's independence assumption. The independence assumption, which presupposes that the residuals are uncorrelated with one another, is essential to the validity of regression analysis.

In conclusion, a Durbin-Watson value of 1.994 indicates that the residuals of the regression model do not substantially display first-order autocorrelation, supporting the independence assumption.

c) Normality Assumption

```
In [56]: # Use qqplot function from StatsModels
fig, ax = plt.subplots()
sm.graphics.qqplot(subset_results_two.resid, dist=stats.norm, line='45', fit=True, ax=ax)

# Customize plot appearance
line = ax.lines[1]
line.set_color("black")
fig.suptitle("Normal Distribution");
```




From the above plot we can see that our model has passed the normality assumption.

d) The Homoscedasticity Assumption

```
In [57]: fig, ax = plt.subplots()

ax.scatter(x = data["price"], y = subset_results_two.resid)
ax.axhline(y=0, color="black")
ax.set_xlabel("price")
ax.set_ylabel("residuals");
```



Model residual plot shows an increase in error rate with increasing house prices

7. Results Interpretation

Based on the correlation matrix, and correlation heatmap, and the model results for the standardised data, squarefoot living variable is the single most significant variable in affecting house prices, then followed by view based on our model.

A 1% increase in sqft_living leads to a 0.84% increase in house price

Since only price, i.e. the target variable, and sqft_living has been been log-transformed, the model interpretation means that for 1 unit increase in floors leads to a 9.4% increase in house prices. A one level improvement in condition leads to a 7.34% increase in house prices while a one level improvement in view leads to a 14% increase house price

8. Limitations and Scope

1. The model's error rate seems to increase with increasing house prices hence predictions for high house prices may not be as accurate as with lower house prices
2. The project's scope is limited to factors directly related to unique property features and the effects of renovations or upgrades. External factors, such as market volatility, regulatory changes, and broader economic shifts, are not included in the model
3. Latitude and Zipcode predictor variables seem to have a significant impact on house prices. However without further context on geography of the regions from which the data emanated it is difficult to interpret this relationship.

9. Recommendations

1. Further to no.3 above, more analysis is required to show how geographical location of a house affect it's price
2. Ways to improve the accuracy, especially for high-value homes, so we can provide even more reliable estimates across the entire range of house prices is needed.
3. Market Segmentation: To determine market segmentation or particular buyer preferences, examine the link between the independent factors and housing prices. For example, a market sector of luxury or high-end homes may be indicated if the prices of houses tend to be higher.

